

Predictive Modelling Birds Classification

Ya Ting Huang
Alexandra Landry
Nastasia Kantsevitch
Julius Alessandro Xanthoudakis



Overview



- Introduction
 - Data Selection
 - Challenges
- Model 1 - Base
- Model 2 - Augmented
- Model 3 - Pre-Trained
- Lessons Learned
- Conclusion

Data Selection

- 15 most represented classes

```
[5] 1 def create_data(df, type_data):
2     processed_data_directory = 'data/'
3     data = []
4     top_15 = df['labels'].value_counts().head(15).index.tolist()
5     for c in top_15:
6         image_folder = os.path.join(processed_data_directory + f'{type_data}/', c)
7         for filename in os.listdir(image_folder):
8             image_path = os.path.join(image_folder, filename)
9             img = imread(image_path, as_gray=False)
10            img = img / 255
11            img = resize(img, (64, 64))
12            class_value = top_15.index(c)
13            data.append([img, class_value])
14     return data
```

```
[6] 1 train_data = create_data(df, 'train')
2     valid_data = create_data(df, 'valid')
3     test_data = create_data(df, 'test')
```

```
[7] 1 print("Training data length: ", len(train_data))
2     print("Validation data length: ", len(valid_data))
3     print("Testing data length: ", len(test_data))
```

```
Training data length: 3306
Validation data length: 75
Testing data length: 75
```

Data Selection cont'd

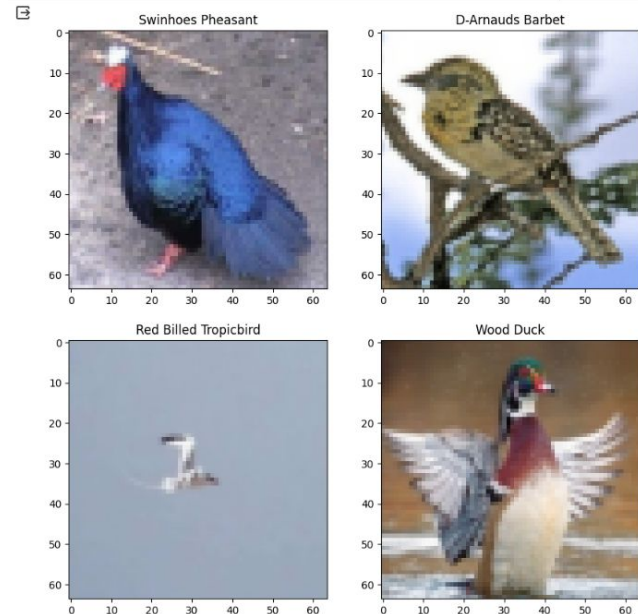
- Retaining Class information

```
[65] 1 def append_label(df, type_data):
      2     top_15 = df['labels'].value_counts().head(15).index.tolist()
      3     label_dict = {}
      4     for idx, c in enumerate(top_15):
      5         label_dict[idx] = c.title()
      6     return label_dict
      7
      8 labels = append_label(df, 'test')
```

1 labels

```
{0: 'Rufous Trepe',
1: 'House Finch',
2: 'D-Arnauds Barbet',
3: 'Ovenbird',
4: 'Asian Green Bee Eater',
5: 'Swinhoes Pheasant',
6: 'Wood Duck',
7: 'Caspian Tern',
8: 'Red Billed Tropicbird',
9: 'Wood Thrush',
10: 'Frill Back Pigeon',
11: 'Pyrrhuloxia',
12: 'Merlin',
13: 'Ornate Hawk Eagle',
14: 'Military Macaw'}
```

```
1 i = 23
2 plt.figure(figsize=(10,10))
3 plt.subplot(221), plt.imshow(X_train[i]), plt.title(labels[int(y_train[i])])
4 plt.subplot(222), plt.imshow(X_train[i+25]), plt.title(labels[int(y_train[i+25])])
5 plt.subplot(223), plt.imshow(X_train[i+50]), plt.title(labels[int(y_train[i+50])])
6 plt.subplot(224), plt.imshow(X_train[i+75]), plt.title(labels[int(y_train[i+75])])
7 plt.show()
8
```



Challenges

- Small Sample Size
- Computational Power
- Vast range of techniques to explore

Model 1

Base

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 222, 222, 32)	896
max_pooling2d (MaxPooling2D)	(None, 111, 111, 32)	0
conv2d_1 (Conv2D)	(None, 109, 109, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 54, 54, 64)	0
conv2d_2 (Conv2D)	(None, 52, 52, 64)	36,928
max_pooling2d_2 (MaxPooling2D)	(None, 26, 26, 64)	0
conv2d_3 (Conv2D)	(None, 24, 24, 128)	73,856
max_pooling2d_3 (MaxPooling2D)	(None, 12, 12, 128)	0
conv2d_4 (Conv2D)	(None, 10, 10, 128)	147,584
max_pooling2d_4 (MaxPooling2D)	(None, 5, 5, 128)	0
conv2d_5 (Conv2D)	(None, 3, 3, 128)	147,584
max_pooling2d_5 (MaxPooling2D)	(None, 1, 1, 128)	0
flatten (Flatten)	(None, 128)	0
dense (Dense)	(None, 256)	33,024
dense_1 (Dense)	(None, 15)	3,855



Total params: 462,223 (1.76 MB)

Trainable params: 462,223 (1.76 MB)

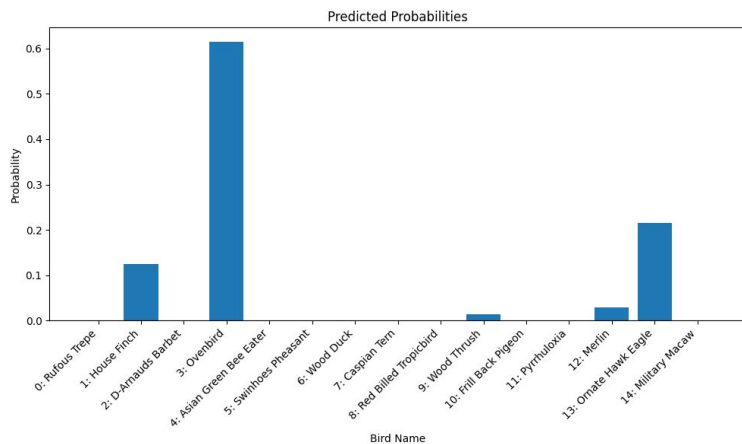

Non-trainable params: 0 (0.00 B)

Epoch 26: early stopping

Restoring model weights from the end of the best epoch: 16.



Test loss: 0.8631170392036438
Test accuracy: 0.8399999737739563



True label: Ornate Hawk Eagle, Predicted label: Ovenbird



Model 2

Augmented

Key Differences

- Image Preprocessing
- More Conv2D Layers
- Batch Normalization
- Regularization within Dense Layer
- Dropout
- Learning Rate

```

1  ### After testing numerous models, this is the best model we could come up with manually.
2  rlrnp = tf.keras.callbacks.ReduceLROnPlateau(monitor='val_loss', factor=0.4, patience=2, verbose=1)
3
4  early_stopping2 = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=4, verbose=1,
5                                                    restore_best_weights=True)
6
7  ## Building the CNN model
8  model2 = keras.Sequential()
9
10 # Data augmentation
11 model2.add(layers.RandomFlip(mode='horizontal_and_vertical', input_shape=(224,224,3)))
12
13 model2.add(Conv2D(32, (3,3), activation='relu'))
14 model2.add(MaxPooling2D((2,2), padding='same'))
15
16 model2.add(Conv2D(64, (3,3), activation='relu'))
17 model2.add(MaxPooling2D((2,2), padding='same'))
18
19 model2.add(Conv2D(128, (3,3), activation='relu'))
20 model2.add(MaxPooling2D((2,2), padding='same'))
21
22 model2.add(Conv2D(256, (3,3), activation='relu'))
23 model2.add(MaxPooling2D((2,2), padding='same'))
24
25 model2.add(Conv2D(256, (3,3), activation='relu'))
26 model2.add(MaxPooling2D((2,2), padding='same'))
27
28 model2.add(layers.BatchNormalization())
29
30 model2.add(Flatten())
31
32 model2.add(layers.Dense(512, kernel_regularizer=regularizers.l2(0.016),
33                        activity_regularizer=regularizers.l1(0.006),
34                        bias_regularizer=regularizers.l1(0.006), activation='relu'))
35
36 model2.add(layers.Dropout(0.20))
37
38 model2.add(Dense(15, activation='softmax'))
39
40 ## Compiling the model
41 model2.compile(loss=SparseCategoricalCrossentropy(from_logits=True), optimizer='adam', metrics=['accuracy'])
42
43 ## Printing the summary of the model
44 model2.summary()
45
46 ## Fitting the model
47 epochs=30
48 history2 = model2.fit(X_train, y_train
49                      ,batch_size=32
50                      ,validation_data=[X_valid, y_valid]
51                      ,epochs=epochs
52                      ,callbacks=[rlrnp, early_stopping2])
53



```

Layer (type)	Output Shape	Param #
random_flip_1 (RandomFlip)	(None, 224, 224, 3)	0
conv2d_11 (Conv2D)	(None, 222, 222, 32)	896
max_pooling2d_11 (MaxPooling2D)	(None, 111, 111, 32)	0
conv2d_12 (Conv2D)	(None, 109, 109, 64)	18,496
max_pooling2d_12 (MaxPooling2D)	(None, 55, 55, 64)	0
conv2d_13 (Conv2D)	(None, 53, 53, 128)	73,856
max_pooling2d_13 (MaxPooling2D)	(None, 27, 27, 128)	0
conv2d_14 (Conv2D)	(None, 25, 25, 256)	295,168
max_pooling2d_14 (MaxPooling2D)	(None, 13, 13, 256)	0
conv2d_15 (Conv2D)	(None, 11, 11, 256)	590,080
max_pooling2d_15 (MaxPooling2D)	(None, 6, 6, 256)	0
batch_normalization_1 (BatchNormalization)	(None, 6, 6, 256)	1,024
flatten_2 (Flatten)	(None, 9216)	0
dense_4 (Dense)	(None, 512)	4,719,104
dropout_1 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 15)	7,695

Total params: 5,706,319 (21.77 MB)


Trainable params: 5,705,807 (21.77 MB)


Non-trainable params: 512 (2.00 KB)



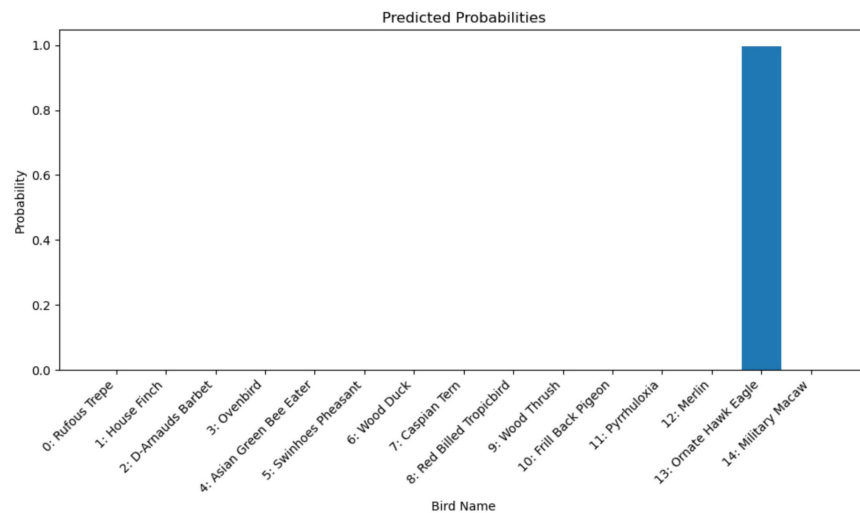
Test loss: 0.498855859041214
Test accuracy: 0.9466666579246521

The learning rate decreased in steps as the epochs progressed





True label: Ornate Hawk Eagle, Predicted label: Ornate Hawk Eagle



Model 3

Pre-Trained

Key Differences

- VGG16 Pre-Trained Model
- Possibility to Freeze/Unfreeze
- No Data Augmentation
- Longer Training Time

```

1  ### Testing pretrained model -- VGG16
2
3  ## Building the CNN model - VGG16 pretrained model
4  base_model_vgg = VGG16(include_top = False, input_shape = (224, 224, 3))
5  base_model_vgg.trainable = False
6
7  # Define the hidden layers and the output
8  x = base_model_vgg.get_layer('block5_pool').output
9  x = layers.GlobalAveragePooling2D()(x)
10 outputs = layers.Dense(15, activation = 'softmax')(x)
11
12 # Create the model
13 model_vgg = models.Model(base_model_vgg.input, outputs)
14
15 ## Compiling the model
16 model_vgg.compile(loss = 'sparse_categorical_crossentropy'
17                   ,optimizer = optimizers.Adam(learning_rate = 0.01)
18                   ,metrics = ['accuracy'])
19
20 ## Printing the summary of the model
21 model_vgg.summary()
22
23 ## Fitting the model
24 epochs=30
25 history_vgg = model_vgg.fit(X_train, y_train
26                             ,epochs = epochs
27                             ,validation_data = [X_valid, y_valid]
28                             ,batch_size=32
29                             ,callbacks=[early_stopping])
30
31 ## Plot for the training results
32 plot_model_training(history_vgg)
33
34 ## Evaluating the model on unseen data with the test data
35 score_vgg = model_vgg.evaluate(X_test, y_test, verbose=0)
36 print("Test loss:", score_vgg[0])
37 print("Test accuracy:", score_vgg[1])
38
39 # Part of the code from: https://medium.com/@bobbycxy/birds-classification-with-pre-trai



```

Layer (type)	Output Shape	Param #
input_layer_3 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1,792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36,928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73,856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147,584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295,168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590,080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590,080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
global_average_pooling2d (GlobalAveragePooling2D)	(None, 512)	0
dense_6 (Dense)	(None, 15)	7,695


Total params: 14,722,383 (56.16 MB)

Trainable params: 7,695 (30.06 KB)

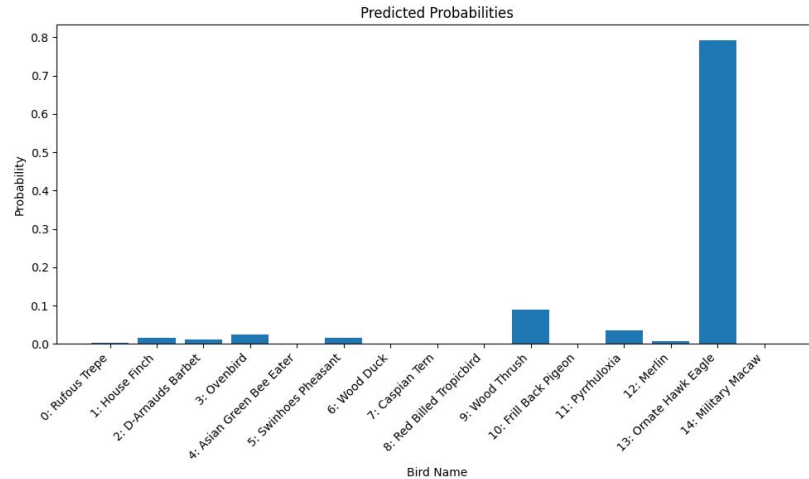
Non-trainable params: 14,714,688 (56.13 MB)



Test loss: 0.13346274197101593
Test accuracy: 0.9599999785423279





True label: Ornate Hawk Eagle, Predicted label: Ornate Hawk Eagle



InterceptV3


- Much quicker to train
- Similar Results to VGG16
- Used Data Augmentation (to counter overfitting)

Test loss: 0.5461175441741943
Test accuracy: 0.8799999952316284




Lessons Learned


- Choose classes more carefully
- Version compatibility
- Utilizing GitHub to manage versions
- Use parameters that are tested and work
- Note on model selection:



Basic



Augmented



VGG16

Conclusion

